

# Python in HPC

Joseph John

National Computational Infrastructure, Australia



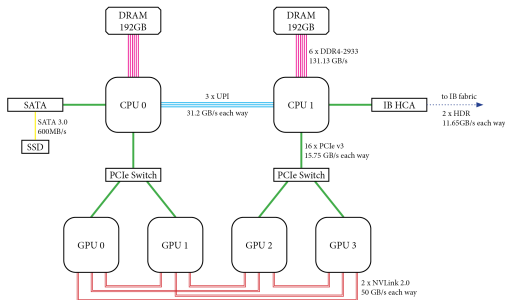
## We will cover:

- Limitation of Python
- Numba
- CuPy
- Dask

- **project:** vp91
- **modules:** python3/3.9.2 cuda/12.0.0
- **venv:** dask-python3.9-venv
- **repo:** <https://github.com/nci900/AAPP-Parallel-Python>

# Why do we need Parallelism?


- Resource utilization is minimal:
  - 1 > 90% compute comes from GPU.
  - 2 CPU threads are under utilized.



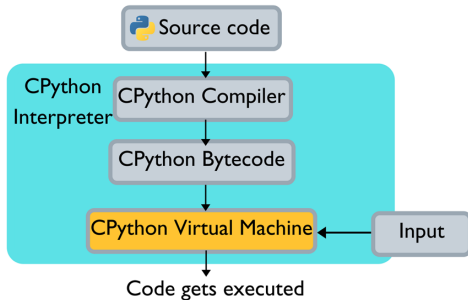
## Limitation of

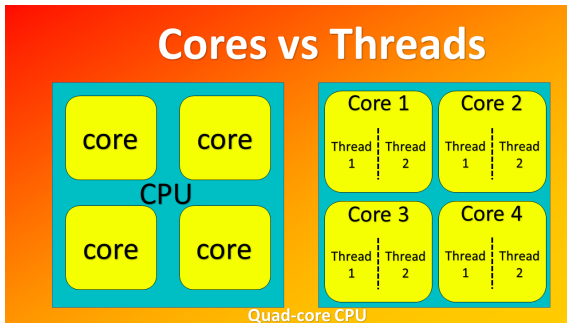


- Python is an interpreted language.
- Interpreter first compiles the source code into platform-independent Bytecode.
- Bytecode is executed by Python Virtual Machine.




- CPython - implements the Python interpreter using both C and Python.
- standard reference implementation of Python.
- Not *thread-safe*.





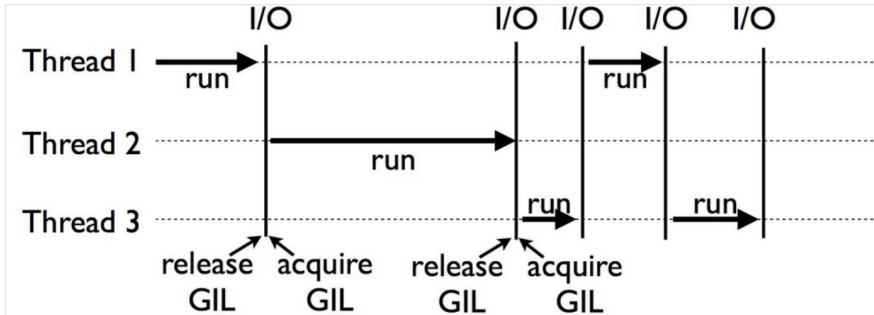


- Memory is shared between multiple threads within a process.
- Single instruction, multiple data. (SIMD)



# Global Interpreter Lock (GIL)?

- Prevent multiple native threads from causing unwanted interactions.
- Only one thread can access Python interpreter at any given time.
- Only one thread to execute the Bytecode at any given time.



# Why do we need GIL?

- Python interpreter is not thread safe.
- Variables are managed a by a refcount.
- Without GIL there can be race conditions.

```
>>> import sys
>>> a = []
>>> b = a
>>> sys.getrefcount(a)
3
```


- **Critical section:** Section of code that accesses shared resources (variables or data structures).
- When multiple threads tries to write to a resource in a critical section it can result in unintended values.
- This is called race condition.
- Reading concurrently does not result in a race condition.

- We prevent race condition using **mutex**.
- Mutex is a synchronization primitive that grants exclusive access to the shared resource to only one thread.
- Only one thread can acquire a mutex at a time.
- Only the thread that acquired the mutex can enter critical section.
- The thread that have access to the mutex should release it (after the critical section), for other threads to acquire it.

# Python's multiprocessing Module

- Spawn multiple native sub-processes within a program.
- Each sub-process is allocated its own memory.

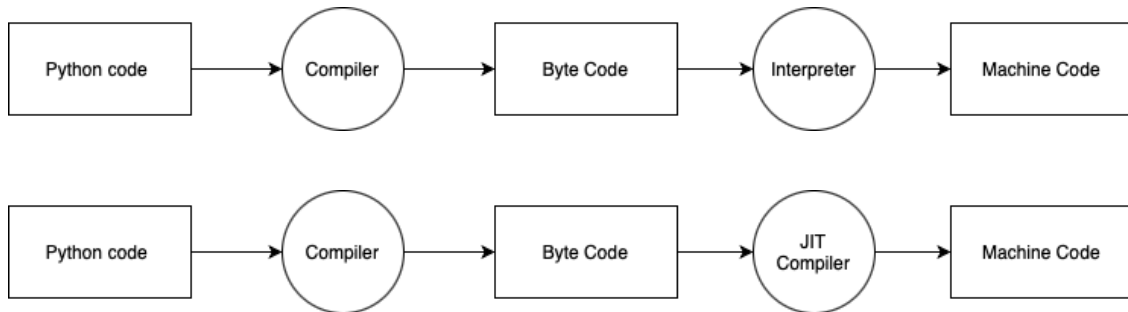
- Each process can run on different CPU cores.
- GIL is not involved.
- Resource intensive when compared to threads.





- 1 Optimized the inefficient use-cases of Numpy
- 2 Multi-dimensional array (ndarray)
- 3 Custom Python C extensions not required





<b>Matrix Size</b>	<b>Numba</b>	<b>C</b>
64 x 64	463x	453x
128 x 128	454x	407x
256 x 256	280x	263x
512 x 512	276x	268x

Lam, Siu Kwan et al. "Numba: a LLVM-based Python JIT compiler." LLVM '15 (2015).

```
def uppercase_decorator(function):  
    def wrapper():  
        func = function()  
        make_uppercase = func.upper()  
        return make_uppercase  
  
    return wrapper
```

```
def say_hi():  
    return 'hello there'
```

```
decorate = uppercase_decorator(say_hi)  
decorate()
```

```
@uppercase_decorator  
def say_hi():  
    return 'hello there'  
  
say_hi()
```

```
@jit  
def f(x, y):  
    return x + y
```

- Decorating the function with `@jit` will mark a function for optimization by Numba's JIT compiler
- The compilation will be deferred until the first function execution
- Different function invocation will result in different compilation

```
@jit(int32(int32, int32))  
def f(x, y):  
    return x + y
```

- We can tell numba to generate code only for one set of arguments

```
@jit
def f(x, y):
    return x + y
```

- Decorating the function with `@jit` will mark a function for optimization by Numba's JIT compiler
- The compilation will be deferred until the first function execution
- Different function invocation will result in different compilation

```
@jit
def square(x):
    return x ** 2
```

```
@jit
def hypot(x, y):
    return math.sqrt(square(x) + square(y))
```

- One compiled function can call another compiled function.



```
x = np.arange(100).reshape(10, 10)
```

```
@jit(nopython=True)
def with_numba(a):
    trace = 0.0
    for i in range(a.shape[0]):
        trace += np.tanh(a[i, i])
    return a + trace
```

- @jit decorator fundamentally operates in two compilation modes, nopython mode and object mode.
- *nopython* compilation mode compile the decorated without the involvement of the Python interpreter.

```
@jit(nogil=True)
def f(x, y):
    return x + y
```

- Release GIL.
- Runs concurrently with other threads executing Python or Numba code.
- Takes advantage of multi-core systems.

```
@jit(cache=True)
def f(x, y):
    return x + y
```

- The chances are you call the same function again and again with the same argument type.
- Numba can cache the compiled code.

```
@jit(nopython=True, parallel=True)
def f(x, y):
    return x + y
```

- This feature only works on CPUs.

```
@vectorize([float64(float64, float64)])  
def sinacosb_vect(a, b):  
    return math.sin(a) * math.cos(b)
```

- Creating a ufunc that operates on a ndarray of a particular type is not straight forward.
- Numba makes this process easy.

## 1 CPU

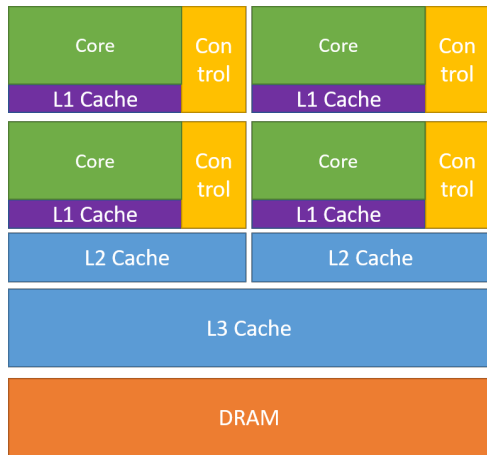
- ▶ Optimized to execute a code as fast as possible
- ▶ Executes a few tens of threads in parallel
- ▶ Transistors are give proportional importance control flow, computation and data caching

## 2 GPU

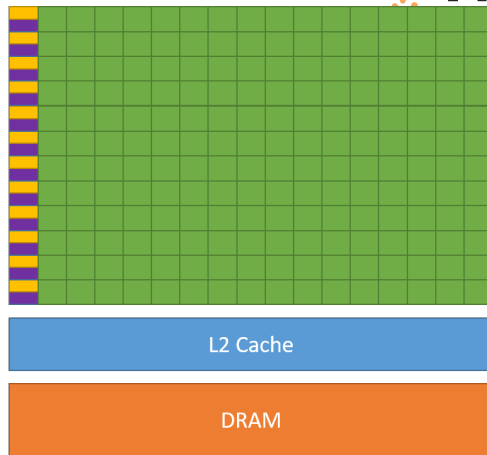
- ▶ Optimized to execute a code as parallel as possible
- ▶ Executes a few thousand of threads in parallel
- ▶ Transistors are disproportional favour computation

<https://docs.nvidia.com/cuda/cuda-c-programming-guide/index.html>

# CPU and GPU

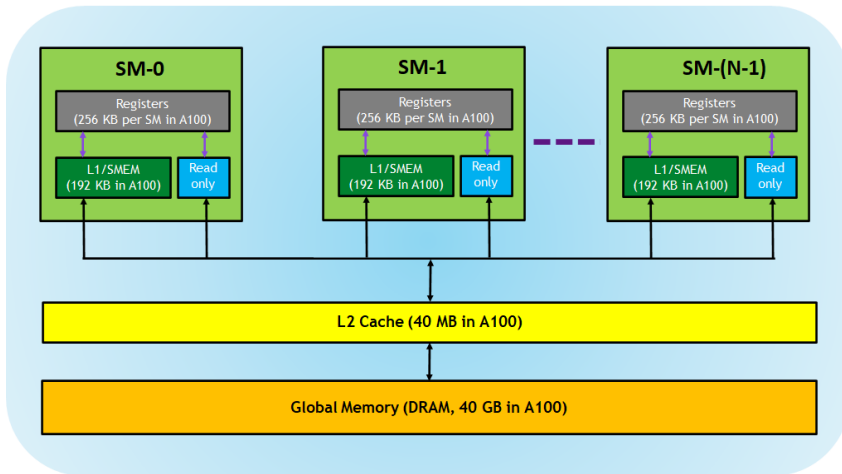


CPU



GPU

<https://docs.nvidia.com/cuda/cuda-c-programming-guide/index.html>




<https://developer.nvidia.com/blog/cuda-refresher-cuda-programming-model/>



- Allocate memory in GPU memory
- Move data from main memory to GPU memory
- Launch GPU kernel
- Move data back to main memory

## Example Program

```
// Allocate Memory  
cudaMalloc(&d_x, N*sizeof(float));  
  
// Move data to GPU  
cudaMemcpy(d_x, x, N*sizeof(float), cudaMemcpyHostToDevice);  
  
// Launch kernel  
increment<<<(N+511)/512, 512>>>(N, d_x);  
  
// Move data back  
cudaMemcpy(x, d_x, N*sizeof(float), cudaMemcpyDeviceToHost)
```



<https://docs.nvidia.com/cuda/cuda-c-programming-guide/index.html>

```
@vectorize(['int64(int64, int64)'], target='cuda')  
def add_ufunc(x, y):  
    return x + y
```

Numba automates the following:

- Allcated GPU memory.
- Copy data to the GPU memory.
- Executed the CUDA kernel with the *correct kernel dimensions given the input sizes*.
- Copy data to the hist memory.
- Return the result as a NumPy array.



CuPy

## 1 CPU

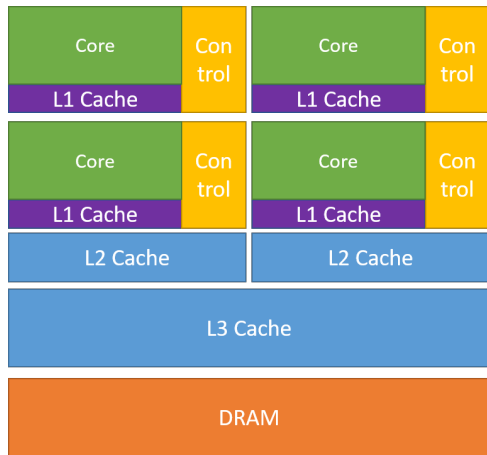
- ▶ Optimized to execute a code as fast as possible
- ▶ Executes a few tens of threads in parallel
- ▶ Transistors are give proportional importance control flow, computation and data caching

## 2 GPU

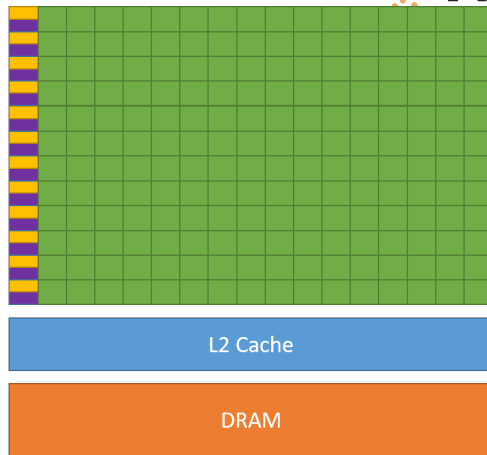
- ▶ Optimized to execute a code as parallel as possible
- ▶ Executes a few thousand of threads in parallel
- ▶ Transistors are disproportional favour computation

<https://docs.nvidia.com/cuda/cuda-c-programming-guide/index.html>

# CPU and GPU



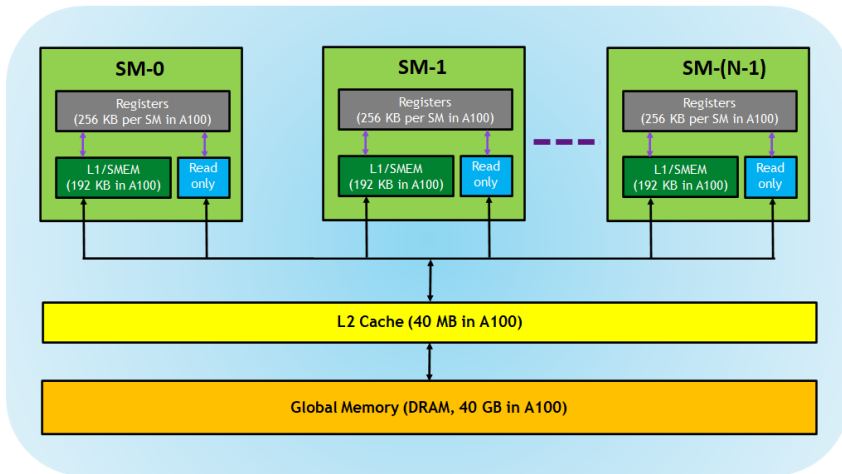
CPU



GPU

<https://docs.nvidia.com/cuda/cuda-c-programming-guide/index.html>





<https://developer.nvidia.com/blog/cuda-refresher-cuda-programming-model/>

- Allocate memory in GPU memory
- Move data from main memory to GPU memory
- Launch GPU kernel
- Move data back to main memory

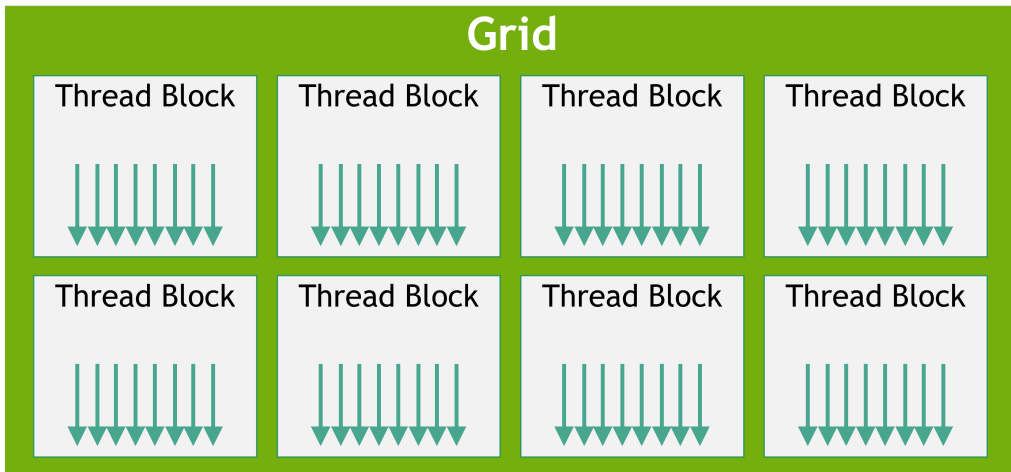
## Example Program

```
// Allocate Memory
cudaMalloc(&d_x, N*sizeof(float));

// Move data to GPU
cudaMemcpy(d_x, x, N*sizeof(float), cudaMemcpyHostToDevice);

// Launch kernel
increment<<<(N+511)/512, 512>>>(N, d_x);

// Move data back
cudaMemcpy(x, d_x, N*sizeof(float), cudaMemcpyDeviceToHost)
```



<https://docs.nvidia.com/cuda/cuda-c-programming-guide/index.html>

```
cp.cuda.runtime.getDevice()
```

- CuPy has a concept of a current device.
- Default GPU device on which on which all operation of related to CuPy takes place.
- Unless specifically mentioned, all operation tasks place in this default device.

```
x_gpu = cp.array([1, 2, 3])
```

- `cupy.array()` allocates the data in the GPU memory.
- If no device is specified the memory gets allocated in the current device.

```
with cp.cuda.Device(1):  
    x_on_gpu1 = cp.array([1, 2, 3, 4, 5])
```

- We can use the device context manager to switch between the devices.

```
x_cpu = np.array([1, 2, 3])  
x_gpu_0 = cp.asarray(x_cpu)
```

- In CuPY the memory allocation and data movement can be done in a single operation.



```
with cp.cuda.Device(1):  
    x_gpu_1 = cp.asarray(x_gpu_0)
```

- D2D transfer is faster than H2D transfer.

```
with cp.cuda.Device(0):  
    x_cpu = cp.asnumpy(x_gpu_0)
```

```
with cp.cuda.Device(1):  
    x_cpu = x_gpu_1.get()
```

- There are two ways to fetch the data from GPU:
  - ▶ `cupy.ndarray.get()`
  - ▶ `cupy.asnumpy()`

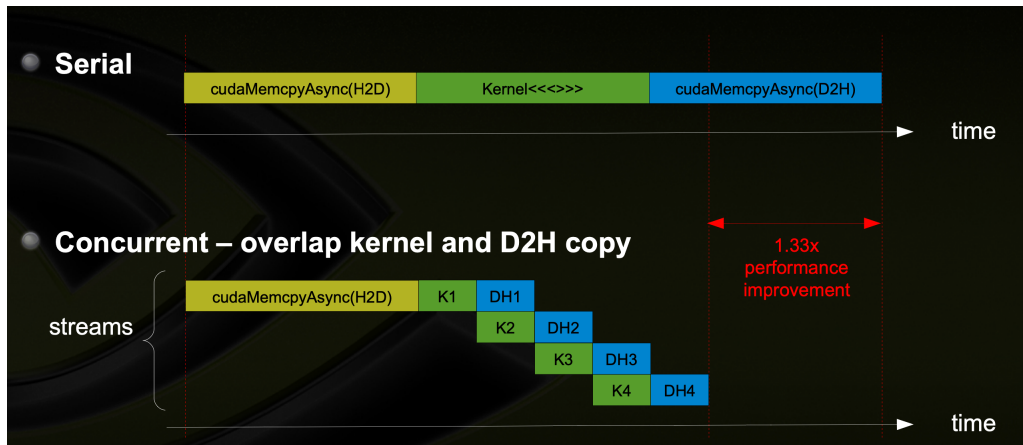
```
def log_array(x):  
    xp = cp.get_array_module(x)  
    xp.log1p(xp.exp(-abs(x)))
```

```
log_array(x_cpu)
```

```
log_array(x_gpu)
```

- We can make function calls to a data, without the knowledge of where they reside.

- Elementwise Kernel.
- Reduction Kernel.
- Raw kernel.



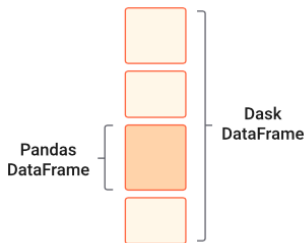
<https://developer.download.nvidia.com/CUDA/training/StreamsAndConcurrencyWebinar.pdf>



- Parallel and distributed computing library for python
- Dask scale up to your full laptop capacity and out to a cloud cluster
- Multi-core and distributed+parallel execution on larger-than-memory datasets

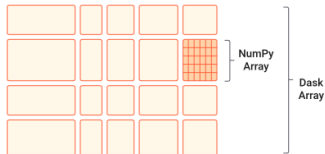
- **High-level collections:** Mimic NumPy, lists, and pandas but can operate in parallel on datasets that don't fit into memory
  - ▶ Array
  - ▶ DataFrame
  - ▶ Bag
- **Low-level collections:** Give finer control to build custom parallel and distributed computations.
  - ▶ Delayed
  - ▶ Futures





- One Dask DataFrame is comprised of many in-memory pandas DataFrames separated along the index
- One operation on a Dask DataFrame triggers many pandas operations on the constituent pandas DataFrames
- These operations are mindful of potential parallelism and memory constraints

- Dask constructs the logic (called task graph) of your computation immediately
- Evaluates them only when necessary



- Dask Array implements a subset of the NumPy ndarray interface using **blocked** algorithms
- Large array is cut into many small arrays
- Large computations are performed by combining many smaller computations

- A Block of code can have operations that can happen in parallel
- Normally in python this would happen sequentially or the user will identify the parallel section and write parallel codes
- The Dask `**delayed**` function decorates your functions so that they operate lazily
- Dask will defer execution of the function, placing the function and its arguments into a task graph
- Dask will then identify opportunities for parallelism in the task graph
- The Dask schedulers will exploit this parallelism, generally improving performance

- We can submit individual functions for evaluation
- The call returns immediately, giving one or more future
  - ▶ whose status begins as “pending”
  - ▶ later becomes “finished”
- There is no **\*\*blocking\*\*** of the local Python session.
- Difference between futures and delayed
  - ▶ delayed is lazy (it just constructs a graph)
  - ▶ futures are eager
- With futures, as soon as the inputs are available and there is compute available, the computation starts

- Dask executes the computations transformation to the distributed data.
- Compute: Converts it to a local object.
- Persist: The object remains distributed.